

FFTnet

March 23, 2019

- 1 We can try to have a neural net 'Learn' the filters, that can be used to tranform X into Y. But as we'll see below, these filters will not be pure sine waves, but will have a lot of noise (with frequencies that are absent in the original signal)

```
In [9]: import numpy as np
import mne
import matplotlib.pyplot as plt

import scipy.io as sio
import tensorflow as tf
config = tf.ConfigProto()
config.gpu_options.allow_growth = True
sess = tf.Session(config=config)
from keras import backend as K
K.set_session(sess)
#####

from keras.layers import Input, Dense, Conv2D, Conv1D, MaxPooling1D, UpSampling1D, Dropout
from keras.layers.advanced_activations import LeakyReLU
from keras import regularizers
# act = LeakyReLU(alpha=0.01)

from keras.models import Model, load_model
from keras.callbacks import TensorBoard, EarlyStopping, Callback
from keras import optimizers
from keras.initializers import Orthogonal as orth

def periodogram(fs, X):
    T = X.shape[0]/fs
    t = np.linspace(0.0, T, int(T/(1/fs)))
    yf = fft(X)
    xf = np.linspace(0.0, fs/2.0, int(len(yf)/2))
    yf = 2/(T*fs) * np.abs(yf[0:int(len(yf)/2)])
    return xf, yf
```

```
In [13]: mat_contents = sio.loadmat('/home/amplifier/home/DATASETS/XY.mat')
```


```
X = mat_contents['X']
Y = mat_contents['Y']

X = X.reshape(1,32,3991,1)
Y = Y.reshape(1,32,3991,1)

fig = plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(np.squeeze(X[:,14,:,:]))
plt.plot(np.squeeze(Y[:,14,:,:]))
plt.title('Time Domain')
plt.legend(['X', 'Y'])

plt.subplot(1,2,2)
xf1, yf1 = periodogram(fs, X[:,14,:,:].flatten())
xf2, yf2 = periodogram(fs, Y[:,14,:,:].flatten())
plt.plot(xf1, yf1)
plt.plot(xf2, yf2)
plt.title('Frequency Domain')
plt.legend(['X', 'Y'])
```

```
Out[13]: <matplotlib.legend.Legend at 0x7f1d0a30fda0>
```



```
In [21]: ker_width = 400
         inpt = Input(shape=(32, 3991, 1))
         l = Conv2D(1, (1,ker_width),
                   use_bias=None,
                   padding="same",
                   activation=None)(inpt)
         model = Model(inpt,l)
         model.compile(loss='binary_crossentropy', optimizer='adam')
```

```

model.summary()

train_history = model.fit(X, Y, epochs=1000, verbose=0)
wts = np.array(model.layers[1].get_weights()).flatten()

fig = plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(wts)
plt.title('Filter in Time Domain')
plt.grid()

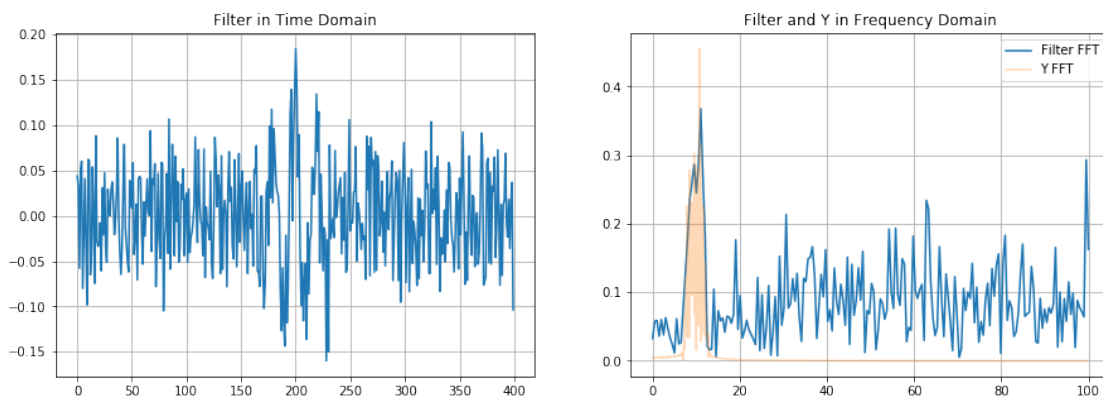
plt.subplot(1,2,2)
plt.title('Filter and Y in Frequency Domain')
xf1, yf1 = periodogram(fs, wts)
plt.plot(xf1, 20*yf1)
xf2, yf2 = periodogram(fs, Y[:,14,:,:].flatten())
plt.plot(xf2, yf2, alpha=1.3)
plt.legend(['Filter FFT', 'Y FFT'])
plt.grid()

```

```

-----
Layer (type)                Output Shape          Param #
=====
input_8 (InputLayer)        (None, 32, 3991, 1)   0
-----
conv2d_8 (Conv2D)           (None, 32, 3991, 1)   400
=====
Total params: 400
Trainable params: 400
Non-trainable params: 0
-----

```



- 2 Alternatively, we can have the model learn NOT the filters directly (when these filters start to include irrelevant frequencies that don't 'resonate' with the frequencies in the input signal, because these aren't there), but have it learn the filter parameter (i.e. frequency of the sinusoid, or Morlet wavelet). In that case we can get pure (and easily interpretable) filters. The results of this filtering can be combined downstream and/or fed into a FC layer for regression or classification. Keras, as far as I know, can't do it, but PyTorch can:

```
In [29]: import torch
         from torch.autograd import Variable
         import torch.optim as optim
         import torch.nn.functional as F
         from scipy import signal
         from scipy.fftpack import fft
         import scipy.io as sio
         import numpy as np
         import matplotlib.pyplot as plt

         def gaussian(x, mu, sig):
             return np.exp(-np.power(x - mu, 2.) / (2 * np.power(sig, 2.)))

         mat_contents = sio.loadmat('/home/amplifier/home/DATASETS/XY.mat')
         X = mat_contents['X']
         Y = mat_contents['Y']

         X = X[:,0:1000]
         Y = Y[:,0:1000]

In [30]: fs = 200
         f = Variable(torch.tensor([2.0, 5.0, 10.0, 15.0, 20.0]), requires_grad=True)
         print(f)
         pi = torch.tensor(np.pi)
         T = 1
         t = torch.tensor(np.linspace(0.0, T, int(T/(1/fs)), dtype='float32'))

         x_values = np.linspace(-4, 4, 200)
         g = torch.tensor(gaussian(x_values, 0, 1)).float()

         k = []
         for i in range(5):
             k.append(g * torch.cos(2.0*pi*f[i]*t))

         fig = plt.figure(figsize=(15,5))
         plt.subplot(1,2,1)
```

```


for i in range(5):
    plt.plot(t.detach().numpy(), k[i].detach().numpy())

t = np.linspace(-T/2, T/2, T * fs, endpoint=False) # sampled @ 200 Hz for 0.5 s
I = []
k_m = []
for j in range(5):
    i, q, e = signal.gausspulse(t, f[j].detach().numpy(), retquad=True, retenv=True)
    I.append(i)
    k_m.append(torch.tensor(i).float())

plt.subplot(1,2,2)
for i in range(5):
    plt.plot(t, I[i])

tensor([ 2.,  5., 10., 15., 20.], requires_grad=True)

```



```

In [31]: def forward(inputs, f, pi):
    fs = 200
    T = 1
    t = torch.tensor(np.linspace(0.0, T, int(T/(1/fs)), dtype='float32'))
    x_values = np.linspace(-4, 4, 200)
    g = torch.tensor(gaussian(x_values, 0, 1)).float()
    k = []
    for i in range(5):
        k.append(g * torch.cos(2.0*pi*f[i]*t))
    filters = torch.stack(k).view(5,1,1,200)
    x = F.conv2d(inputs, filters, padding=(0,padding))
    return x, filters

In [32]: signal_len = X.shape[1]
    n_eeg_chans = X.shape[0]
    ker_width = 200

```

```

padding = 99
n_filters = 5
electrode = 14

inputs = torch.tensor(X).view(1, 1, n_eeg_chans, signal_len) # batch, in_channel

# GET FILTERED SIGNAL AND FILTERS BASED ON SET FREQUENCIES AND USING MY "MORLET" THAT S
x, filters = forward(inputs, f, pi)

# GET FILTERED SIGNAL AND FILTERS BASED ON SET FREQUENCIES AND USING MORLET THAT DOESN
filters_m = torch.stack(k_m).view(5,1,1,200)
x_m = F.conv1d(inputs, filters_m, padding=(0,padding))


# GET FILTERED SIGNAL USING NUMPY'S CONV. JUST TO SEE THAT IT WORKS THE SAME AS TORCH.M
x_c = []
for i in range(n_filters):
    x_c.append(np.convolve(X[electrode,:], filters_m[i,0,0,:].detach().numpy().flatten(
x_c = np.vstack(x_c)

print('inputs.shape =\t', inputs.shape)
print('filters.shape =\t' , filters.shape)
print('filters_m.shape =\t' , filters_m.shape)
print('x.shape =\t', x.shape)
print('x_m.shape =\t', x_m.shape)
print('x_c.shape =\t', x_c.shape)

fig = plt.figure(figsize=(15,10))
for band in range(0, filters.shape[0]):
    plt.subplot(filters.shape[0]+1, 1, band+1)
    plt.plot(X[electrode,:])
    plt.plot(x[0,band,electrode,:].detach().numpy())
    plt.plot(x_m[0,band,electrode,:].detach().numpy())
    plt.plot(x_c[band])
    plt.legend(['original', 'conv2d, plain cos', 'conv2d, morlet', 'conv, morlet'])
    plt.title('Electrode: ' + str(electrode) + ', Band: ' + str(band))

inputs.shape =          torch.Size([1, 1, 32, 1000])
filters.shape =          torch.Size([5, 1, 1, 200])
filters_m.shape =        torch.Size([5, 1, 1, 200])
x.shape =                torch.Size([1, 5, 32, 999])
x_m.shape =              torch.Size([1, 5, 32, 999])
x_c.shape =              (5, 1000)

```




```
In [33]: fs = 200.0
         xf1, yf1 = periodogram(fs, X[electrode,:])
         xf2, yf2 = periodogram(fs, Y[electrode,:])

         fig = plt.figure(figsize=(15,5))

         plt.subplot(1,2,1)
         plt.plot(X[electrode,:])
         plt.plot(Y[electrode,:])

         plt.subplot(1,2,2)
         plt.plot(xf1, yf1)
         plt.plot(xf2, yf2)
         plt.grid()
```



```

In [45]: optimizer = optim.Adam([f]) # tell the optimizer which var we want optimized
        criterion = torch.nn.MSELoss() # set loss type

        x, filters = forward(inputs, f, pi) # forward the model to get filtered signal and filt

        predictions = torch.sum(x, 1).view(32,-1)
        targets = torch.tensor(Y[:,0:999]).float()
        loss = criterion(predictions, targets) # compute loss

        print('predictions: \trequires_grad?\t', predictions.requires_grad)
        print('targets: \trequires_grad?\t', targets.requires_grad)
        print('loss: \t\trequires_grad?\t', loss.requires_grad)

predictions:          requires_grad?          True
targets:              requires_grad?          False
loss:                  requires_grad?          True

In [44]: for epoch in range(5000):
        x, filters = forward(inputs, f, pi)
        predictions = torch.sum(x, 1).view(32,-1)
        loss = criterion(predictions, targets)
        loss.backward(retain_graph=True)
        optimizer.step() # update the variable being optimized (f)
        optimizer.zero_grad()
        if epoch % 100 == 0:
            print('loss:\t', np.round(loss.item(), 1), '\tfreqs:\t', f)

loss:      8164.0      freqs:      tensor([ 2.0010,  4.9990, 10.0010, 15.0010, 19.9990])
loss:      8115.3      freqs:      tensor([ 2.0558,  4.9871, 10.0521, 15.0697, 19.9084])
loss:      8080.7      freqs:      tensor([ 2.1268,  5.0512, 10.0597, 15.0460, 19.8152])
loss:      8014.2      freqs:      tensor([ 2.2251,  5.1467, 10.0657, 14.9872, 19.6710])
loss:      7952.5      freqs:      tensor([ 2.3136,  5.2366, 10.0706, 14.9217, 19.5093])
loss:      7910.7      freqs:      tensor([ 2.3792,  5.3039, 10.0713, 14.8631, 19.3543])
loss:      7882.6      freqs:      tensor([ 2.4243,  5.3497, 10.0687, 14.8162, 19.2097])
loss:      7860.5      freqs:      tensor([ 2.4543,  5.3798, 10.0673, 14.7886, 19.0714])
loss:      7832.6      freqs:      tensor([ 2.4749,  5.4010, 10.0841, 14.8265, 18.9190])
loss:      7376.0      freqs:      tensor([ 2.5026,  5.4328, 10.2078, 15.2148, 18.6048])
loss:      6906.4      freqs:      tensor([ 2.5501,  5.4818, 10.2731, 15.4376, 18.3694])
loss:      6872.7      freqs:      tensor([ 2.5943,  5.5257, 10.2933, 15.4099, 18.3027])
loss:      6847.4      freqs:      tensor([ 2.6335,  5.5639, 10.3042, 15.3609, 18.2511])
loss:      6826.1      freqs:      tensor([ 2.6700,  5.5989, 10.3108, 15.3145, 18.2019])
loss:      6808.4      freqs:      tensor([ 2.7057,  5.6325, 10.3155, 15.2720, 18.1561])
loss:      6793.9      freqs:      tensor([ 2.7426,  5.6670, 10.3198, 15.2338, 18.1144])
loss:      6782.0      freqs:      tensor([ 2.7830,  5.7045, 10.3256, 15.2007, 18.0775])
loss:      6771.5      freqs:      tensor([ 2.8301,  5.7483, 10.3348, 15.1731, 18.0460])

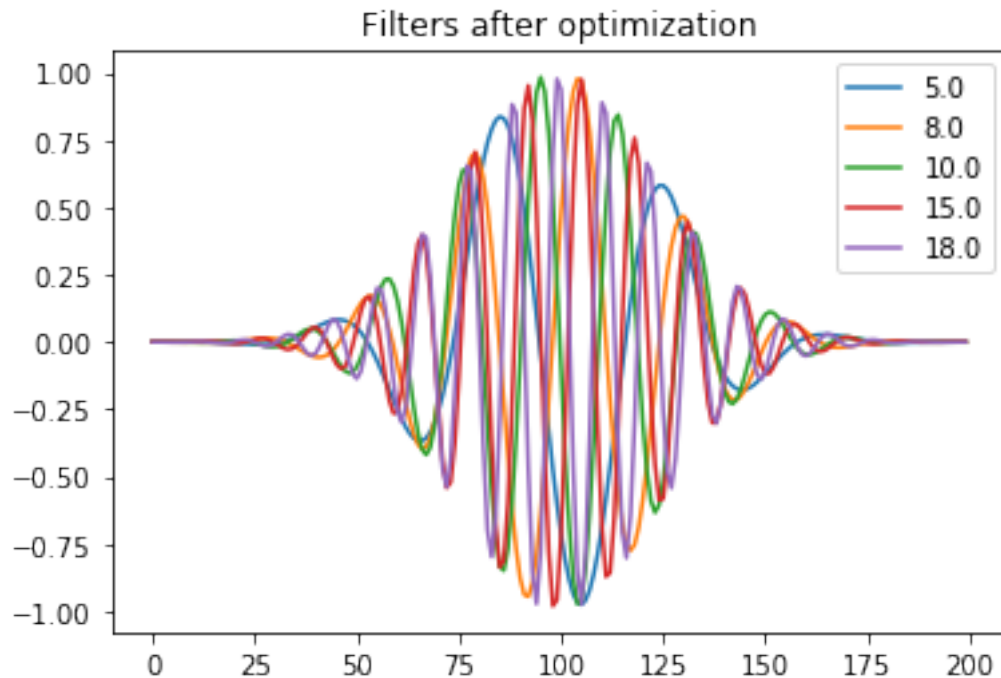
```


loss:	6761.1	freqs:	tensor([2.8877, 5.8017, 10.3492, 15.1516, 18.0204]
loss:	6749.1	freqs:	tensor([2.9583, 5.8678, 10.3703, 15.1369, 18.0014]
loss:	6734.8	freqs:	tensor([3.0403, 5.9454, 10.3977, 15.1295, 17.9897]
loss:	6720.1	freqs:	tensor([3.1258, 6.0274, 10.4280, 15.1294, 17.9854]
loss:	6707.5	freqs:	tensor([3.2056, 6.1049, 10.4568, 15.1350, 17.9879]
loss:	6697.8	freqs:	tensor([3.2749, 6.1730, 10.4809, 15.1445, 17.9953]
loss:	6690.7	freqs:	tensor([3.3346, 6.2321, 10.4989, 15.1556, 18.0054]
loss:	6684.9	freqs:	tensor([3.3884, 6.2858, 10.5065, 15.1659, 18.0159]
loss:	6678.1	freqs:	tensor([3.4466, 6.3460, 10.4606, 15.1683, 18.0226]
loss:	5824.6	freqs:	tensor([3.7322, 6.6993, 9.8938, 15.0856, 17.9837]
loss:	5418.9	freqs:	tensor([3.9839, 6.8959, 9.8263, 14.9938, 17.8926]
loss:	5294.1	freqs:	tensor([4.1002, 7.0098, 9.9298, 14.9244, 17.8160]
loss:	5218.9	freqs:	tensor([4.1961, 7.1026, 10.0085, 14.8854, 17.7663]
loss:	5176.6	freqs:	tensor([4.2719, 7.1757, 10.0691, 14.8732, 17.7452]
loss:	5151.1	freqs:	tensor([4.3319, 7.2338, 10.1180, 14.8797, 17.7459]
loss:	5133.8	freqs:	tensor([4.3809, 7.2815, 10.1592, 14.8973, 17.7599]
loss:	5121.0	freqs:	tensor([4.4222, 7.3218, 10.1951, 14.9204, 17.7810]
loss:	5111.0	freqs:	tensor([4.4581, 7.3570, 10.2270, 14.9454, 17.8049]
loss:	5103.1	freqs:	tensor([4.4901, 7.3884, 10.2558, 14.9704, 17.8291]
loss:	5096.8	freqs:	tensor([4.5189, 7.4166, 10.2821, 14.9942, 17.8523]
loss:	5091.7	freqs:	tensor([4.5451, 7.4423, 10.3062, 15.0163, 17.8741]
loss:	5087.7	freqs:	tensor([4.5691, 7.4657, 10.3284, 15.0367, 17.8940]
loss:	5084.4	freqs:	tensor([4.5910, 7.4871, 10.3487, 15.0553, 17.9122]
loss:	5081.9	freqs:	tensor([4.6111, 7.5067, 10.3675, 15.0722, 17.9288]
loss:	5079.8	freqs:	tensor([4.6295, 7.5246, 10.3847, 15.0876, 17.9438]
loss:	5078.2	freqs:	tensor([4.6464, 7.5409, 10.4005, 15.1015, 17.9574]
loss:	5076.9	freqs:	tensor([4.6618, 7.5559, 10.4150, 15.1142, 17.9697]
loss:	5075.9	freqs:	tensor([4.6758, 7.5695, 10.4283, 15.1257, 17.9808]
loss:	5075.1	freqs:	tensor([4.6886, 7.5819, 10.4405, 15.1361, 17.9909]
loss:	5074.5	freqs:	tensor([4.7003, 7.5932, 10.4515, 15.1455, 18.0000]
loss:	5074.0	freqs:	tensor([4.7109, 7.6035, 10.4616, 15.1540, 18.0082]
loss:	5073.6	freqs:	tensor([4.7205, 7.6128, 10.4708, 15.1617, 18.0156]

```
In [65]: fs = 200
         pi = torch.tensor(np.pi)
         _, filters = forward(inputs, f, pi)

         for i in range(5):
             filt = filters[i,:,:,:].detach().numpy().flatten()
             plt.plot(filt)
         plt.legend(f.detach().numpy().round(decimals=0).tolist())
         plt.title('Filters after optimization')
```

```
Out[65]: Text(0.5, 1.0, 'Filters after optimization')
```



```
In [73]: Y.shape
```

```
Out[73]: (32, 1000)
```

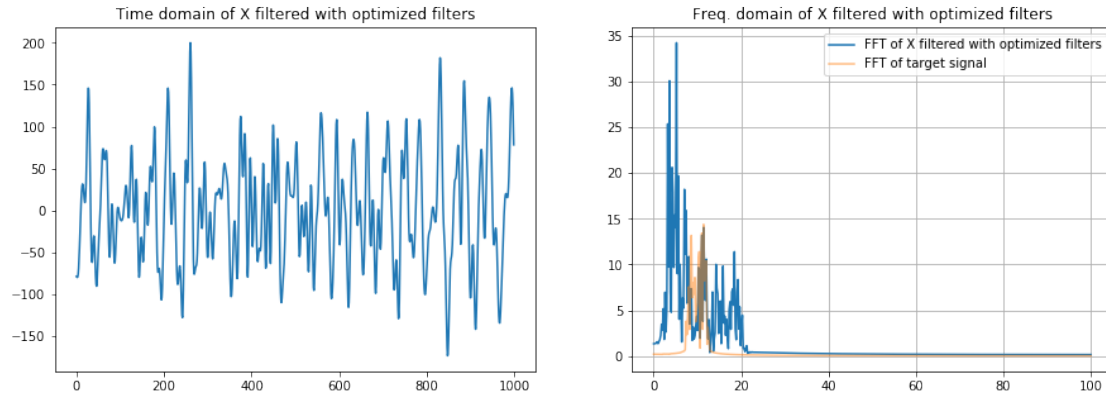
```
In [ ]:
```

```
In [78]: x_ = x[:, :, electrode, :].sum(1).detach().numpy().flatten()
```

```
fig = plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(x_)
plt.title('Time domain of X filtered with optimized filters')

plt.subplot(1,2,2)
xf1, yf1 = periodogram(fs, x_)
plt.plot(xf1, yf1)
xf2, yf2 = periodogram(fs, Y[electrode, :].flatten())
plt.plot(xf2, 30*yf2, alpha=0.5)
plt.grid()
plt.title('Freq. domain of X filtered with optimized filters')
plt.legend(['FFT of X filtered with optimized filters', 'FFT of target signal'])
```

```
Out[78]: <matplotlib.legend.Legend at 0x7f1bec0f66a0>
```



- 3 As we can see, this approach can't fit the target signal very well, but perhaps, if we add weights to each kernel (perhaps an FC layer) we would be able to achieve optimal filters

In []: